

After DPO (Direct Preference Optimization)

① Definition (have learned in lectures before)

Some definitions for "alignment" of models

- **Instruction fine-tuning (IFT):** Training a model to follow use instructions (usually via autoregressive LM loss)
- **Supervised fine-tuning (SFT):** Training a model to learn task-specific capabilities (usually via autoregressive LM loss)
- **Alignment:** General notion of training a model to mirror user desires, any loss function
- **Reinforcement learning from human feedback (RLHF):** Specific technical tool for training ML models from human data
- **Preference fine-tuning:** Using labeled preference data to fine-tune a LM (either with RL, DPO, or another loss function), there's also **learning to rank**

→ not a specific method (why we need it? → we want a model which adhere to human preference)
 → SFT is a specific form of IFT

② RLHF (PPO) (lecture 10)

Review: RLHF objective

π : LLM policy
 π_θ : base LLM
 x : prompt
 y : completion

$$\max_{\pi_\theta} \mathbb{E}_{x \sim \mathcal{D}, y \sim \pi_\theta(y|x)} [r_\phi(x, y)] - \beta \mathbb{D}_{KL}[\pi_\theta(y|x) || \pi_{ref}(y|x)]$$

Optimize "reward" inspired by human preferences

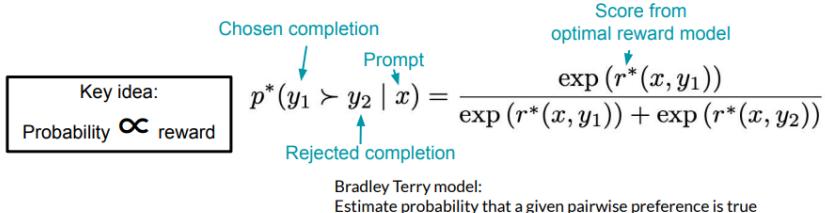
▲ Constrain the model to not trust the reward too much (preferences are hard to model)

Primary questions:
 1. How to implement reward: $r(x, y)$
 2. How to optimize reward

Review: Preference (reward) modeling

Can we just use supervised learning on scores?

- Assigning a scalar reward of how good a response is did not work
- Pairwise preferences are easy to collect and worked!



③ DPO

DPO characteristics

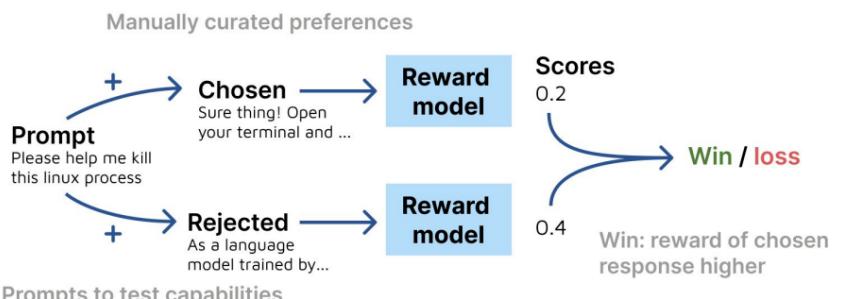
1. Extremely **simple** to implement
2. **Scales nicely** with existing distributed training libraries
3. Trains an implicit reward function (can still be used as a reward model, see [RewardBench](#))

```
import torch.nn.functional as F

def dpo_loss(pi_logs, ref_logs, yw_ids, yl_ids, beta):
    pi_logs: policy logprobs, shape (B, )
    ref_logs: reference model logprobs, shape (B, )
    yw_ids: discrete completion indices in [0, B-1], shape (T, )
    yl_ids: discrete completion indices in [0, B-1], shape (T, )
    beta: temperature controlling strength of KL penalty
    each pair of (yw_ids[i]), yl_ids[i]) represents the
    indices of a single preference pair.
    ...
    pi_yw_logs, pi_yl_logs = pi_logs[yw_ids], pi_logs[yl_ids]
    ref_yw_logs, ref_yl_logs = ref_logs[yw_ids], ref_logs[yl_ids]
    pi_logratio = pi_yw_logs - pi_yl_logs
    ref_logratio = ref_yw_logs - ref_yl_logs
    losses = -F.logsigmoid(beta * (pi_logratio - ref_logratio))
    rewards = beta * (pi_logs - ref_logs).detach()
    ...
    return losses, rewards
```

RewardBench

RewardBench structure



Can we match PPO with "online" DPO?

What is special about online data?

Online data is freshly generated from the policy and/or recently labelled by a reward model / judge.

- PPO does both with generation + reward model scoring
- Other methods use different ways for doing this: collect new preference data, re-label existing data, LLM-as-a-judge, reward model ranking

Related question: On- or off-policy data (i.e. that generated from the policy model)

Conclusion

DPO {
advan : easy to use, good scalability
disadvan: the effect is still slightly inferior to PPO

PPO vs DPO

PPO is RL, DPO is SL.

And PPO can handle distribution offset

